

Optimization for Machine Learning: Gradient Descent and Stochastic Gradient Descent

Mathematical Foundations of Data Science

Difan Zou

Nov. 11, 2025

Recap: Optimization Basics (Previous Lecture)

Reference: Lecture_note_optimization_basics.pdf

Key concepts covered:

- **Optimization problem formulation:** $\min_{\theta} f(\theta)$ subject to constraints
- **Convexity:** A function is convex if $f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y)$
- **Optimality conditions:**
 - Unconstrained: $\nabla f(\theta^*) = 0$
 - Constrained: KKT conditions
- **Smoothness (Lipschitz gradient):** $\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$
- **Strong convexity:** $f(y) \geq f(x) + \nabla f(x)^T(y - x) + \frac{\mu}{2}\|y - x\|^2$

Today's lecture: We apply these concepts to **Machine Learning optimization**

- Empirical Risk Minimization (ERM)
- Gradient Descent and its convergence analysis
- Stochastic Gradient Descent for large-scale problems

Outline

- 1 Introduction to Optimization in ML
- 2 Gradient Descent
- 3 Examples: Convex Loss Functions
- 4 Limitations and Challenges
- 5 Stochastic Gradient Descent
- 6 Application: GD for Linear Regression
- 7 Backpropagation for Neural Networks
- 8 Summary and Takeaways

Empirical Risk Minimization (ERM)

Goal: Given training data points, find parameter θ that minimizes loss

Optimization Problem:

$$\theta_S^* := \arg \min_{\theta \in \Theta} L_S(h_\theta)$$

Model Training Objective:

$$L_S(h_\theta) = \frac{1}{n} \sum_{i=1}^n \ell(f_\theta(x_i), y_i)$$

where:

- n : training sample size
- ℓ : loss function
- f_θ : model function with parameter θ
- (x_i, y_i) : training data (feature, target)

When Do Exact Solutions Exist?

Example 1: Linear Regression (Ordinary Least Squares)

- **Objective:** $\min_{\theta} \frac{1}{2n} \sum_{i=1}^n (\theta^\top x_i - y_i)^2$
- **Closed-form solution:** $\theta^* = (X^\top X)^{-1} X^\top y$

Example 2: Ridge Regression

- **Objective:** $\min_{\theta} \|X\theta - y\|^2 + \lambda \|\theta\|^2$
- **Closed-form solution:** $\theta^* = (X^\top X + \lambda I)^{-1} X^\top y$

Example 3: Linear Programming

- **Form:** $\min_x c^\top x$ subject to $Ax \leq b, x \geq 0$
- **Solvers:** Simplex method, interior point methods (polynomial time)

Linear Regression: Closed-Form Solution Derivation (Part 1)

Objective: $\min_{\theta} L(\theta) = \frac{1}{2n} \sum_{i=1}^n (\theta^\top x_i - y_i)^2$

Matrix notation: Let $X \in \mathbb{R}^{n \times d}$ be the data matrix (rows are samples), $y \in \mathbb{R}^n$ be the targets

$$L(\theta) = \frac{1}{2n} \|X\theta - y\|^2$$

Expand the squared norm:

$$\begin{aligned} L(\theta) &= \frac{1}{2n} (X\theta - y)^\top (X\theta - y) \\ &= \frac{1}{2n} (\theta^\top X^\top X\theta - 2\theta^\top X^\top y + y^\top y) \end{aligned}$$

Goal: Find θ^* such that $\nabla_{\theta} L(\theta^*) = 0$

Linear Regression: Closed-Form Solution Derivation (Part 2)

Compute the gradient: Starting from

$$L(\theta) = \frac{1}{2n}(\theta^\top X^\top X \theta - 2\theta^\top X^\top y + y^\top y)$$

Gradient rules:

- $\nabla_\theta(\theta^\top A \theta) = 2A\theta$ for symmetric A
- $\nabla_\theta(\theta^\top b) = b$
- $\nabla_\theta(c) = 0$ for constant c

Apply gradient:

$$\begin{aligned}\nabla_\theta L(\theta) &= \frac{1}{2n}(2X^\top X \theta - 2X^\top y) \\ &= \frac{1}{n}(X^\top X \theta - X^\top y)\end{aligned}$$

Linear Regression: Closed-Form Solution Derivation (Part 3)

Set gradient to zero:

$$\nabla_{\theta} L(\theta) = \frac{1}{n}(X^T X \theta - X^T y) = 0$$

Solve for θ :

$$X^T X \theta - X^T y = 0$$

$$X^T X \theta = X^T y$$

Multiply both sides by $(X^T X)^{-1}$:

$$\theta^* = (X^T X)^{-1} X^T y$$

Note: This requires $X^T X$ to be invertible (full column rank)

- If $d > n$ (more features than samples): not invertible
- If features are linearly dependent: not invertible
- Solution: regularization (Ridge regression)

Ridge Regression: Closed-Form Solution Derivation (Part 1)

Objective: $\min_{\theta} L(\theta) = \frac{1}{2n} \|X\theta - y\|^2 + \frac{\lambda}{2} \|\theta\|^2$

Expand:

$$\begin{aligned}L(\theta) &= \frac{1}{2n} (X\theta - y)^\top (X\theta - y) + \frac{\lambda}{2} \theta^\top \theta \\&= \frac{1}{2n} (\theta^\top X^\top X\theta - 2\theta^\top X^\top y + y^\top y) + \frac{\lambda}{2} \theta^\top \theta\end{aligned}$$

Compute gradient:

$$\begin{aligned}\nabla_{\theta} L(\theta) &= \frac{1}{n} (X^\top X\theta - X^\top y) + \lambda\theta \\&= \frac{1}{n} X^\top X\theta + \lambda\theta - \frac{1}{n} X^\top y\end{aligned}$$

Ridge Regression: Closed-Form Solution Derivation (Part 2)

Set gradient to zero:

$$\frac{1}{n}X^T X \theta + \lambda \theta - \frac{1}{n}X^T y = 0$$

Rearrange:

$$\begin{aligned}\frac{1}{n}X^T X \theta + \lambda \theta &= \frac{1}{n}X^T y \\ \frac{1}{n}(X^T X + n\lambda I)\theta &= \frac{1}{n}X^T y \\ (X^T X + n\lambda I)\theta &= X^T y\end{aligned}$$

Solve for θ :

$$\theta^* = (X^T X + n\lambda I)^{-1} X^T y$$

Key property: $X^T X + n\lambda I$ is always invertible for $\lambda > 0$

- Works even when $d > n$
- Regularization ensures numerical stability

Why We Still Need Iterative Methods: Computational Complexity

Linear Regression Closed-Form:

- Computing $(X^\top X)^{-1}$: $O(d^3)$ time for d -dimensional features
- Matrix-vector products: $O(nd^2)$ time
- **Total:** $O(nd^2 + d^3)$

Problem: For large d (e.g., $d = 10^6$ in modern ML):

- $d^3 \approx 10^{18}$ operations → **computationally infeasible!**

Gradient Descent Alternative:

- Per iteration: $O(nd)$ to compute gradient
- K iterations: $O(Knd)$ total
- If $K \ll d$, much faster than $O(d^3)$

Example: $K = 100, d = 10^6, n = 10^4$

- GD: $\approx 10^{12}$ operations
- Closed-form: $\approx 10^{18}$ operations (**million times slower!**)

Why We Still Need Iterative Methods: Memory Constraints

Closed-form solution requires:

- Storing $X^\top X$: $d \times d$ matrix $\rightarrow O(d^2)$ memory
- For $d = 10^6$: need $\sim 8\text{TB}$ of memory (assuming 8 bytes per float)

Iterative methods require:

- Store parameters θ : $O(d)$ memory
- Store one mini-batch: $O(Bd)$ memory ($B \ll n$)
- **Huge difference:** GB vs TB

Why We Still Need Iterative Methods: No Closed-Form Solution

Many ML models have **no closed-form solution**:

Logistic Regression:

- Objective: $\min_{\theta} \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y_i \theta^\top x_i))$
- Issue: No closed-form due to nonlinear log/exp
- Solution: Must use iterative methods (GD, Newton's method)

Neural Networks:

- Objective: $\min_{\theta} \frac{1}{n} \sum_{i=1}^n \ell(f_{\theta}(x_i), y_i)$
- f_{θ} : multi-layer composition with nonlinear activations
- Issue: Highly non-convex, no closed-form solution
- Solution: SGD with backpropagation is the standard

Comparison of Optimization Methods

Method	Time	Memory	Streaming	General
Closed-form (OLS)	$O(nd^2 + d^3)$	$O(d^2)$	✗	Linear only
Gradient Descent	$O(Knd)$	$O(d)$	✗	Most models
Linear Programming	Polynomial	Varies	✗	Linear only

K : number of iterations

Note: Other iterative methods (e.g., stochastic methods) will be discussed later

High-Dimensional Parameter Spaces:

- Modern neural networks: millions to billions of parameters
 - ResNet-50: ~25 million parameters
 - GPT-3: 175 billion parameters
 - BERT-Large: 340 million parameters
- Image classification: input dimension $d = 224 \times 224 \times 3 \approx 150,000$
- **No hope** for closed-form solutions or traditional methods

Extremely Complicated Objectives

Consider a deep neural network loss:

$$L(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(f_\theta(x_i), y_i)$$

where

$$f_\theta(x) = W^{(L)}\sigma(W^{(L-1)}\sigma(\dots\sigma(W^{(1)}x + b^{(1)})\dots) + b^{(L)})$$

Properties:

- **Highly non-convex:** countless local minima, saddle points, plateaus
- **No analytical form:** composition of many nonlinear functions
- **No structure** amenable to traditional solvers (LP, QP, SDP)
- **Black-box nature:** can only evaluate function and gradient

Numerical Methods: The ONLY Practical Choice

For modern empirical risk minimization:

- ✗ No closed-form solutions
- ✗ No polynomial-time exact algorithms known
- ✗ Traditional optimization (LP, convex opt.) doesn't apply
- ✓ **Gradient-based methods** are the only scalable approach
- ✓ Can leverage **automatic differentiation** (backpropagation)
- ✓ Works despite non-convexity (empirically successful!)

Bottom Line

For deep learning and large-scale ML, gradient descent and its variants (especially SGD) are not just convenient—they are the **only viable optimization approach** we currently have.

Algorithm Basics

Core Idea: Starting from a point, move in the negative gradient direction

Update Rule:

$$\theta^{(t+1)} = \theta^{(t)} - \eta_t \cdot \nabla_{\theta} L(\theta^{(t)})$$

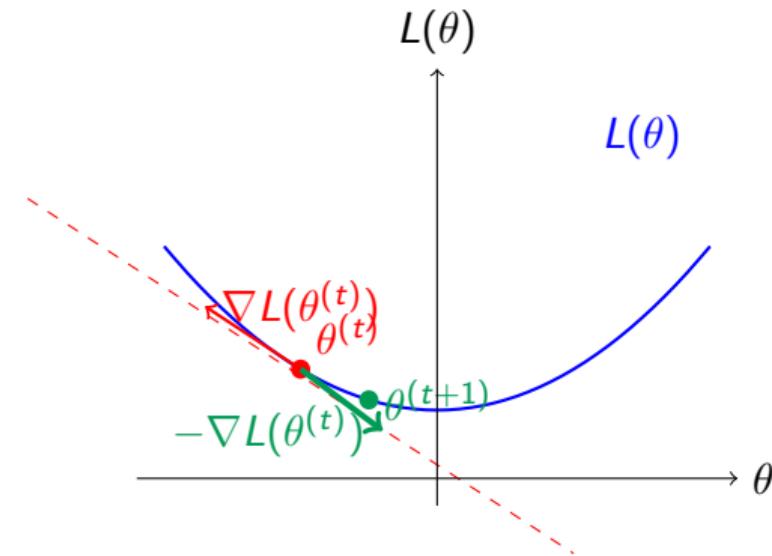
where:

- η_t : step size (learning rate)
- $\nabla_{\theta} L(\theta^{(t)})$: gradient at iteration t

Intuition:

- Linear approximation: $L(\theta + \Delta) \approx L(\theta) + \nabla_{\theta} L(\theta) \cdot \Delta$
- Best descent direction is negative gradient: $-\nabla_{\theta} L(\theta)$

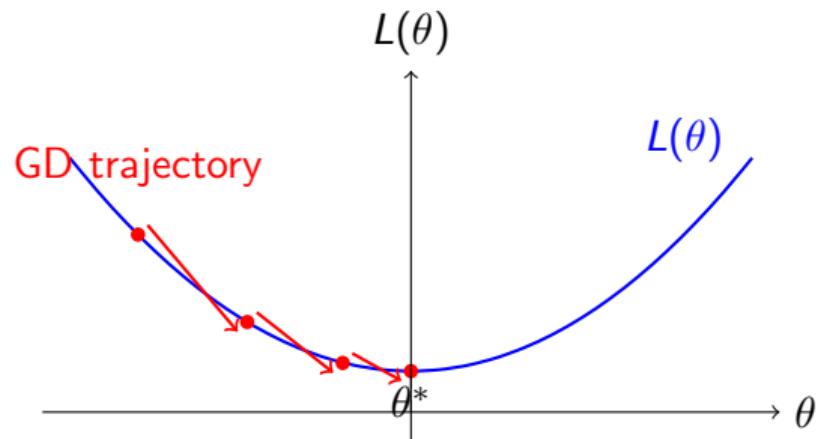
Gradient Descent: Geometric Intuition



Key insight: Negative gradient points in the direction of steepest descent

Outcome of Gradient Descent

Outcome: With proper stepsize, GD finds a point with zero gradient



Only if $\nabla_{\theta} L(\theta)$ approaches zero, the objective can stop decreasing

Smoothness Assumption

Definition: Function is L -smooth if $\lambda_{\max}(\nabla^2 L(\theta)) \leq L$ for all θ

Convergence with Smoothness:

By smoothness:

$$L(\theta + \Delta) \leq L(\theta) + \nabla L(\theta) \cdot \Delta + \frac{L}{2} \|\Delta\|^2$$

Setting $\Delta = -\eta \nabla L(\theta)$:

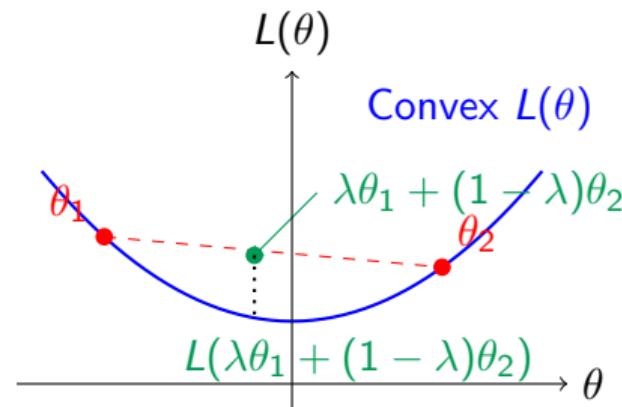
$$\begin{aligned} L(\theta^{(t+1)}) &\leq L(\theta^{(t)}) - \eta \|\nabla L(\theta^{(t)})\|^2 + \frac{L\eta^2}{2} \|\nabla L(\theta^{(t)})\|^2 \\ &= L(\theta^{(t)}) - (\eta - L\eta^2/2) \|\nabla L(\theta^{(t)})\|^2 \end{aligned}$$

Stepsize requirement: $\eta \leq 2/L$ ensures descent, then **Objective keeps decreasing until zero gradient.**

Convex Functions

Definition: For any θ_1, θ_2 and $\lambda \in [0, 1]$:

$$L(\lambda\theta_1 + (1 - \lambda)\theta_2) \leq \lambda L(\theta_1) + (1 - \lambda)L(\theta_2)$$



Key Property: Zero gradient \Rightarrow global minimum in convex case

Convex Functions: Equivalent Definitions

There are multiple equivalent definitions for convex functions:

① **Monotone gradient:**

$$\langle \nabla L(\theta_1) - \nabla L(\theta_2), \theta_1 - \theta_2 \rangle \geq 0$$

② **First-order condition:**

$$L(\theta_1) - L(\theta_2) \geq \langle \nabla L(\theta_2), \theta_1 - \theta_2 \rangle$$

③ **Hessian condition:**

$$\nabla^2 L(\theta) \succeq 0 \text{ (positive semi-definite)}$$

If θ^* satisfies $\nabla L(\theta^*) = 0$, then $L(\theta) - L(\theta^*) \geq 0$ for any $\theta \Rightarrow L(\theta^*)$ is the global minimum

Examples of Convex Loss Functions

Example 1: Linear Regression

- Loss: $\ell(f(x), y) = \frac{1}{2}(\theta^\top x - y)^2$
- Gradient: $\nabla_\theta \ell = (\theta^\top x - y)x$
- Empirical risk: $L_S(\theta) = \frac{1}{2n} \sum_{i=1}^n (\theta^\top x_i - y_i)^2$
- Gradient: $\nabla_\theta L_S = \frac{1}{n} \sum_{i=1}^n (\theta^\top x_i - y_i)x_i$
- **Property:** Convex, can be strongly convex with n data points

Example 2: Logistic Regression

- Loss: $\ell(f(x), y) = \log(1 + \exp(-y \cdot \theta^\top x))$
- Gradient: $\nabla_\theta \ell = -y \cdot x / (1 + \exp(y \cdot \theta^\top x))$
- **Property:** Convex

Convergence Rate: Smooth Objective

Theorem

Assume the objective function is L -smooth. Let $\eta = 1/L$, we have

$$\|\nabla L(\theta^{(t)})\|^2 = O(1/t)$$

Proof Sketch:

- ① From smoothness with $\eta = 1/L$:

$$L(\theta^{(t+1)}) \leq L(\theta^{(t)}) - \frac{1}{2L} \|\nabla L(\theta^{(t)})\|^2$$

- ② Sum over iterations:

$$\sum_{k=0}^t \|\nabla L(\theta^{(k)})\|^2 \leq 2L[L(\theta^{(0)}) - L(\theta^{(t+1)})]$$

- ③ Average $\leq O(1/t) \Rightarrow$ best $\leq O(1/t)$

Convergence Rate: Smooth + Convex

Theorem

Assume the objective function is L -smooth and convex. Let $\eta = 1/L$, we have

$$L(\theta^{(t)}) - L(\theta^*) = O(1/t)$$

Key Steps:

- ① By convexity: $L(\theta^{(t+1)}) - L(\theta^*) \leq \langle \nabla L(\theta^{(t)}), \theta^{(t)} - \theta^* \rangle - \frac{\eta}{2} \|\nabla L(\theta^{(t)})\|^2$
- ② Using update rule:
$$\|\theta^{(t+1)} - \theta^*\|^2 = \|\theta^{(t)} - \theta^*\|^2 - 2\eta \langle \nabla L(\theta^{(t)}), \theta^{(t)} - \theta^* \rangle + \eta^2 \|\nabla L(\theta^{(t)})\|^2$$
- ③ Combine: $L(\theta^{(t+1)}) - L(\theta^*) \leq \frac{1}{2\eta} [\|\theta^{(t)} - \theta^*\|^2 - \|\theta^{(t+1)} - \theta^*\|^2]$
- ④ Sum over iterations: telescoping sum gives $O(1/t)$

Strongly Convex Functions

Definition: Function is μ -strongly convex if:

- $\langle \nabla L(\theta_1) - \nabla L(\theta_2), \theta_1 - \theta_2 \rangle \geq \mu \|\theta_1 - \theta_2\|^2$
- Hessian: $\nabla^2 L(\theta) \succeq \mu I$

Theorem

Assume L -smooth and μ -strongly convex. Let $\eta = 1/L$ and $\kappa = L/\mu$, we have

$$L(\theta^{(t)}) - L(\theta^*), \|\theta^{(t)} - \theta^*\| = O(\exp(-t/\kappa))$$

- Exponential convergence (much faster than convex case!)
- κ is the condition number

Summary of Convergence Rates

Function Property	Guarantee	Rate
L -smooth	$\ \nabla L(\theta^{(t)})\ ^2 \leq \epsilon$	$O(1/t)$
L -smooth + convex	$L(\theta^{(t)}) - L(\theta^*) \leq \epsilon$	$O(1/t)$
L -smooth + μ -strongly convex	$L(\theta^{(t)}) - L(\theta^*) \leq \epsilon$	$O(\exp(-t/\kappa))$

- Stronger assumptions \Rightarrow faster convergence
- Strong convexity gives exponential (linear in log scale) convergence
- Condition number $\kappa = L/\mu$ determines convergence speed

Experimental Verification: Convergence Rates

Setup: Run gradient descent on different function types and observe convergence

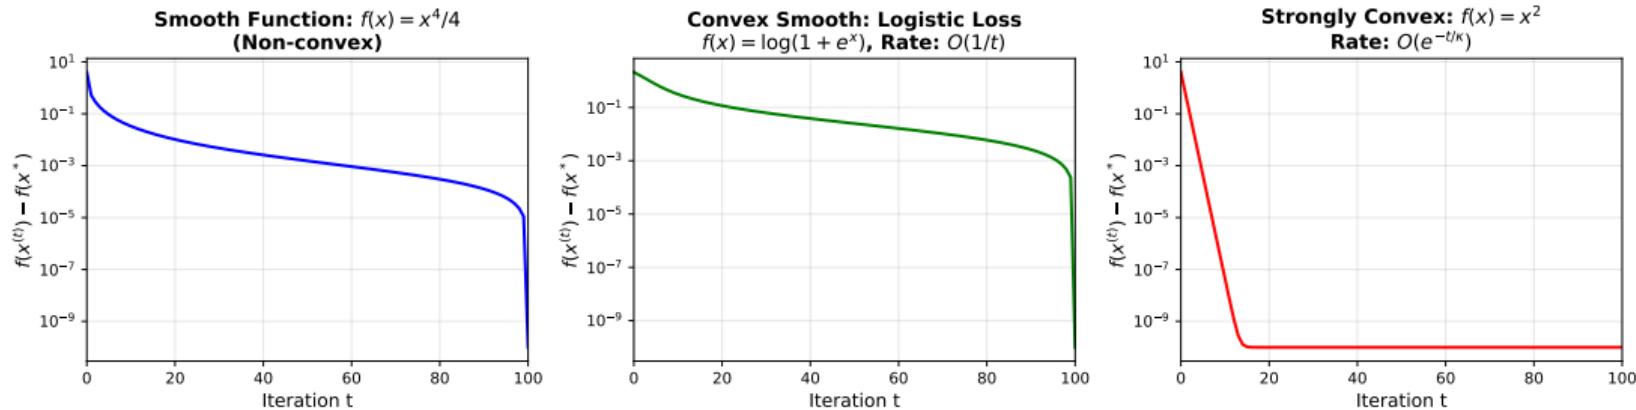
Functions tested:

- ① **Smooth (non-convex):** $f(x) = x^4/4$
- ② **Convex smooth:** $f(x) = \log(1 + e^x)$ (logistic loss)
- ③ **Strongly convex:** $f(x) = x^2$

Experimental protocol:

- Initial point: $x^{(0)} = 2.0$
- Run GD for 100 iterations with appropriate learning rates
- Track $f(x^{(t)}) - f(x^*)$ (optimality gap)
- Plot on log scale to see convergence rate

Convergence Plots: Different Function Classes



Observations:

- Strongly convex: **exponential decay** (straight line on log plot)
- Convex smooth: **$O(1/t)$ decay** (slower, sublinear)
- Smooth (non-convex): Eventually reaches local minimum

Linear Regression: Detailed Gradient Derivation

Goal: Compute $\nabla_{\theta} \ell(f(x), y)$ where $\ell(f(x), y) = \frac{1}{2}(\theta^T x - y)^2$

Step 1: Apply chain rule. Let $u = \theta^T x - y$, then:

$$\frac{\partial \ell}{\partial \theta} = \frac{\partial \ell}{\partial u} \cdot \frac{\partial u}{\partial \theta}$$

Step 2: Compute $\frac{\partial \ell}{\partial u}$:

$$\ell = \frac{1}{2}u^2 \Rightarrow \frac{\partial \ell}{\partial u} = u = \theta^T x - y$$

Step 3: Compute $\frac{\partial u}{\partial \theta}$ where $u = \theta^T x - y$:

$$\frac{\partial u}{\partial \theta} = \frac{\partial(\theta^T x)}{\partial \theta} = x$$

Step 4: Combine using chain rule:

$$\nabla_{\theta} \ell = (\theta^T x - y) \cdot x$$

Logistic Regression: Detailed Gradient Derivation

Goal: Compute $\nabla_{\theta} \ell(f(x), y)$ where $\ell = \log(1 + \exp(-y \cdot \theta^T x))$

Step 1: Apply chain rule. Let $u = -y \cdot \theta^T x$, then:

$$\frac{\partial \ell}{\partial \theta} = \frac{\partial}{\partial \theta} \log(1 + e^u) = \frac{1}{1 + e^u} \cdot \frac{\partial}{\partial \theta} (1 + e^u)$$

Step 2: Compute $\frac{\partial}{\partial \theta} (1 + e^u)$:

$$\frac{\partial}{\partial \theta} (1 + e^u) = e^u \cdot \frac{\partial u}{\partial \theta} = e^{-y \cdot \theta^T x} \cdot (-y \cdot x)$$

Step 3: Combine:

$$\begin{aligned}\nabla_{\theta} \ell &= \frac{e^{-y \cdot \theta^T x}}{1 + e^{-y \cdot \theta^T x}} \cdot (-y \cdot x) \\ &= \frac{-y \cdot x}{1 + e^{y \cdot \theta^T x}}\end{aligned}$$

Logistic Regression: Alternative Form of Gradient

Rewriting the gradient: Starting from

$$\nabla_{\theta} \ell = \frac{-y \cdot x}{1 + e^{y \cdot \theta^T x}}$$

Alternative formulation: Multiply numerator and denominator by $e^{-y \cdot \theta^T x}$:

$$\begin{aligned}\nabla_{\theta} \ell &= \frac{-y \cdot x \cdot e^{-y \cdot \theta^T x}}{e^{-y \cdot \theta^T x} (1 + e^{y \cdot \theta^T x})} \\ &= \frac{-y \cdot x \cdot e^{-y \cdot \theta^T x}}{e^{-y \cdot \theta^T x} + 1}\end{aligned}$$

Sigmoid interpretation: Define $\sigma(z) = \frac{1}{1+e^{-z}}$, then:

$$\nabla_{\theta} \ell = -y \cdot x \cdot (1 - \sigma(y \cdot \theta^T x))$$

Intuition: Gradient proportional to prediction error weighted by input features

Proving Convexity: Linear Regression Loss

Loss function: $\ell(\theta) = \frac{1}{2}(\theta^\top x - y)^2$

Goal: Prove ℓ is convex by showing Hessian $\nabla^2\ell(\theta) \succeq 0$

Step 1: Compute gradient (already done):

$$\nabla_\theta \ell = (\theta^\top x - y) \cdot x$$

Step 2: Compute Hessian (second derivative):

$$\begin{aligned}\nabla_\theta^2 \ell &= \frac{\partial}{\partial \theta} [(\theta^\top x - y) \cdot x] \\ &= \frac{\partial}{\partial \theta} [x \cdot \theta^\top x - y \cdot x] \\ &= x \cdot x^\top\end{aligned}$$

Step 3: Show $\nabla^2\ell = xx^\top$ is positive semi-definite:

- For any vector v : $v^\top (xx^\top)v = (x^\top v)^2 \geq 0 \checkmark$
- Therefore ℓ is convex!

Proving Convexity: Logistic Regression Loss

We can similarly calculate the Hessian:

$$\nabla_{\theta}^2 \ell = \underbrace{\sigma(y \cdot \theta^\top x)(1 - \sigma(y \cdot \theta^\top x))}_{\alpha(\theta) > 0} \cdot xx^\top$$

Key observation: $\alpha(\theta) = \sigma(z)(1 - \sigma(z)) > 0$ for all z because:

- $\sigma(z) \in (0, 1)$ for all $z \in \mathbb{R}$
- Therefore $\sigma(z)(1 - \sigma(z)) > 0$

Positive semi-definiteness: For any $v \in \mathbb{R}^d$:

$$v^\top \nabla^2 \ell \cdot v = \alpha(\theta) \cdot v^\top (xx^\top)v = \alpha(\theta) \cdot (x^\top v)^2 \geq 0$$

Conclusion: Logistic regression loss is convex!

Non-Convex Function Examples

Example 1: Simple non-convex function

$$f(\theta) = \theta^4 - \theta^2, \quad \theta \in \mathbb{R}$$

Analysis:

- Gradient: $f'(\theta) = 4\theta^3 - 2\theta = 2\theta(2\theta^2 - 1)$
- Critical points: $\theta = 0, \pm \frac{1}{\sqrt{2}}$
- Hessian: $f''(\theta) = 12\theta^2 - 2$
- At $\theta = 0$: $f''(0) = -2 < 0$ (**saddle point!**)
- At $\theta = \pm \frac{1}{\sqrt{2}}$: $f''(\pm \frac{1}{\sqrt{2}}) = 4 > 0$ (local minima)

Example 2: Matrix factorization loss

$$L(W, H) = \|X - WH\|_F^2$$

- Non-convex in joint (W, H) despite being convex in each separately
- Used in collaborative filtering, topic modeling

Non-Convex Optimization: The Real Challenge

Reality Check: Most ML problems are **non-convex**

- Neural networks: highly non-convex loss landscapes
- Matrix factorization, dictionary learning: non-convex
- Many clustering objectives: non-convex

Previous Analysis Assumed Convexity:

Our nice convergence guarantees ($O(1/t)$, exponential) only apply to convex functions!

Challenge

In non-convex settings, gradient descent may not find the global minimum

Definition: Point θ^* is a local minimum if:

$$L(\theta^*) \leq L(\theta) \text{ for all } \theta \text{ in a neighborhood of } \theta^*$$

Issue:

- GD converges to a local minimum, not necessarily global
- No gradient information can distinguish local from global at stationary point
- Different initializations → different local minima

Saddle Points

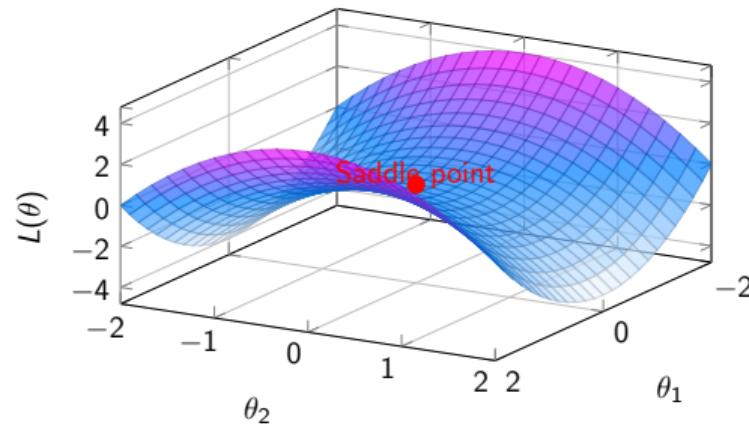
Definition: Point θ^* is a saddle point if:

- $\nabla L(\theta^*) = 0$ (gradient is zero)
- Hessian $\nabla^2 L(\theta^*)$ has both positive and negative eigenvalues
- Not a local minimum or maximum

Property:

- Function decreases in some directions, increases in others
- GD can get **stuck** at saddle points (especially with exact gradient)
- In high dimensions, saddle points are **extremely common**

Saddle Point: Visualization



Example: $L(\theta_1, \theta_2) = \theta_1^2 - \theta_2^2$ has saddle at origin

Mathematical Example: Saddle Point

Simple saddle point:

$$L(\theta_1, \theta_2) = \theta_1^2 - \theta_2^2$$

Gradient:

$$\nabla L = \begin{bmatrix} 2\theta_1 \\ -2\theta_2 \end{bmatrix} = 0 \text{ at origin}$$

Hessian at origin:

$$\nabla^2 L = \begin{bmatrix} 2 & 0 \\ 0 & -2 \end{bmatrix}$$

- Positive eigenvalue (2): upward curvature in θ_1 direction
- Negative eigenvalue (-2): downward curvature in θ_2 direction
- Origin is a saddle point, **NOT a minimum**

High-Dimensional Saddle Points

Key Insight (Dauphin et al., 2014): In high dimensions:

- Saddle points are exponentially more common than local minima
- Most critical points ($\nabla L = 0$) are saddle points, not local minima

Why? For random function in d dimensions:

- Probability that all d eigenvalues of Hessian are positive: $(1/2)^d \rightarrow 0$ as $d \rightarrow \infty$
- Local minimum requires ALL eigenvalues positive
- Saddle point needs just ONE negative eigenvalue: much more likely

Implication for Deep Learning:

- High-dimensional parameter spaces (millions of parameters)
- GD rarely gets trapped at *bad* local minima
- More likely to encounter saddle points
- Fortunately: SGD noise helps escape saddle points!

Plateaus and Slow Convergence

Plateau: Region where $\|\nabla L(\theta)\| \approx 0$ but not exactly zero

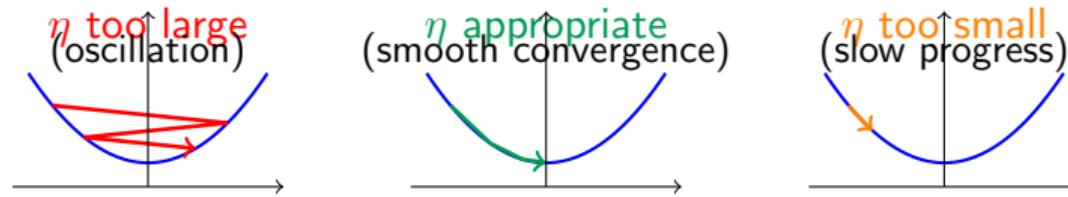
- GD makes very slow progress
- Can take many iterations to escape
- Looks like convergence, but haven't reached optimum

Example in Neural Networks:

- Saturated neurons (sigmoid/tanh): gradients near 0
- Vanishing gradients in deep networks

Sensitivity to Learning Rate

Problem: Step size η critically affects convergence



Challenge:

- Optimal η depends on local curvature (L -smoothness)
- L may vary across parameter space
- Requires tuning or adaptive methods

Summary of Limitations

Challenge	GD Impact	Mitigation
Local minima	Converges to local	Multiple inits
Saddle points	Can get stuck	SGD, add noise
Plateaus	Slow convergence	Adaptive LR
Learning rate	Critical tuning	Grid search

Despite These Limitations

GD/SGD remain the **workhorses** of modern ML

- Empirically successful for deep learning
- Many enhancements exist (momentum, adaptive methods)
- Theory still catching up to practice!

Motivation: Problem with Full-Batch GD

Gradient computation:

$$\nabla L(\theta) = \frac{1}{n} \sum_{i=1}^n \nabla L_i(\theta)$$

- Requires querying **all n training data points**
- Per-iteration complexity: $O(n)$
- **Expensive for large datasets!**

Example:

- ImageNet: $n \approx 1.2$ million images
- Each GD iteration requires forward/backward pass on all images
- For large models, this can take hours per iteration

Incremental Gradient Descent (SGD)

Algorithm:

- ① Randomly permute training data $(x_1, y_1), \dots, (x_n, y_n)$
- ② In iteration t , use single data point (x_t, y_t) :

$$\theta^{(t+1)} = \theta^{(t)} - \eta \cdot \nabla L_t(\theta^{(t)})$$

Advantage: Per-iteration complexity $O(1)$ instead of $O(n)$

Trade-off

- + Much faster per iteration
- Noisier gradient estimates
- May need more iterations to converge

Mini-batch SGD

Algorithm:

- ① Randomly sample mini-batch $\mathcal{I}^{(t)}$ with $|\mathcal{I}^{(t)}| = B$
- ② Update using mini-batch gradient:

$$\theta^{(t+1)} = \theta^{(t)} - \eta \cdot \frac{1}{B} \sum_{i \in \mathcal{I}^{(t)}} \nabla L_i(\theta^{(t)})$$

Properties:

- Expected update equals full-batch gradient:

$$\mathbb{E}[\theta^{(t+1)} | \theta^{(t)}] = \theta^{(t)} - \eta \cdot \nabla L(\theta^{(t)})$$

- Per-iteration complexity: $O(B)$, where $B \ll n$
- Trade-off between computation and variance

Behavior of SGD

Expected Direction: Same as GD in expectation

$$\mathbb{E}[\text{mini-batch gradient}] = \text{full gradient}$$

Variance: Stochastic gradient has variance

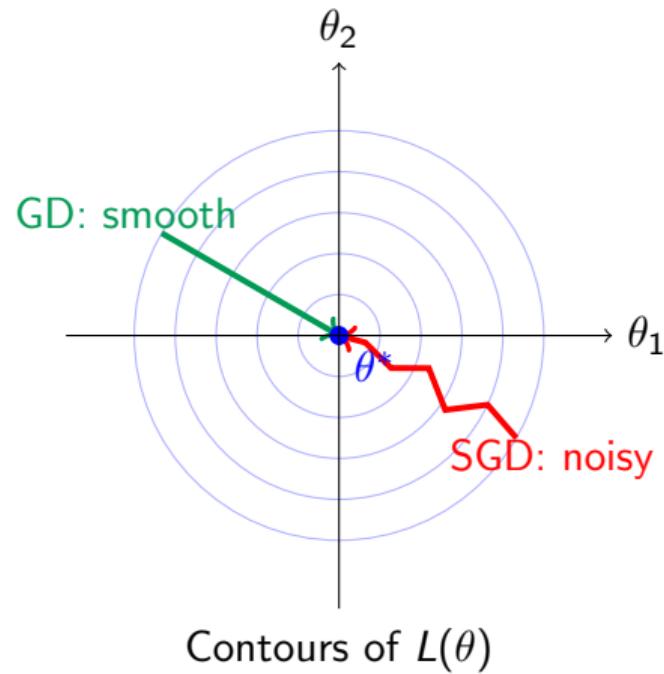
- Large variance: slower convergence, oscillation
- Small variance: behavior similar to GD

Key Insight

When data points are similar:

- Progress of one-step SGD \approx progress of one-step GD
- Complexity of one-step SGD \ll complexity of one-step GD

SGD vs GD: Visual Comparison



SGD takes noisier path but can escape sharp minima

Convergence Analysis: Setup

Assume: All $L_i(\theta)$ are L -smooth

Let g_t be stochastic gradient at iteration t

By smoothness:

$$L(\theta^{(t+1)}) \leq L(\theta^{(t)}) - \eta_t \langle g_t, \nabla L(\theta^{(t)}) \rangle + \frac{L\eta_t^2}{2} \|g_t\|^2$$

Taking expectation:

$$\begin{aligned}\mathbb{E}[L(\theta^{(t+1)})] &\leq L(\theta^{(t)}) - \eta_t \|\nabla L(\theta^{(t)})\|^2 + \frac{L\eta_t^2}{2} \mathbb{E}[\|g_t\|^2] \\ &= L(\theta^{(t)}) - (\eta_t - L\eta_t^2/2) \|\nabla L(\theta^{(t)})\|^2 \\ &\quad + \frac{L\eta_t^2}{2} \mathbb{E}[\|g_t - \nabla L(\theta^{(t)})\|^2]\end{aligned}$$

Convergence Analysis: Two Terms

$$\mathbb{E}[L(\theta^{(t+1)})] \leq L(\theta^{(t)}) - \underbrace{(\eta_t - L\eta_t^2/2)\|\nabla L(\theta^{(t)})\|^2}_{\text{Decreasing term}} + \underbrace{\frac{L\eta_t^2}{2}\mathbb{E}[\|g_t - \nabla L(\theta^{(t)})\|^2]}_{\text{Variance term}}$$

Bounded Variance Assumption:

$$\mathbb{E}[\|g - \nabla L(\theta)\|^2] \leq \sigma^2$$

Key Observations:

- Need small stepsize ($\eta_t \rightarrow 0$) for convergence
- Otherwise SGD oscillates around stationary point
- Zero total gradient \neq zero stochastic gradient

Convergence with Constant Step size

With constant stepsize $\eta \leq 1/L$:

$$\sum_{k=0}^t \eta_k \mathbb{E}[\|\nabla L(\theta^{(k)})\|^2] \leq 2[L(\theta^{(0)}) - \mathbb{E}[L(\theta^{(t)})]] + L \cdot \sum_{k=0}^t \eta_k^2 \sigma^2$$

Implications:

- Constant stepsize: SGD does not converge exactly
- Oscillates in a neighborhood of the optimum
- Neighborhood size depends on η and σ^2

For exact convergence:

- Need decaying stepsize $\eta_t \rightarrow 0$
- E.g., $\eta_t = O(1/t)$ or $\eta_t = O(1/\sqrt{t})$
- But this slows down convergence

SGD vs GD: Sample Complexity Comparison

Question: How many samples do SGD and GD need to achieve ϵ -accuracy?

Gradient Descent (GD):

- Uses all n samples per iteration
- Per-iteration cost: $O(nd)$
- Iterations needed: $O(\log(1/\epsilon))$ for strongly convex, $O(1/\epsilon)$ for convex
- **Total sample accesses:** $O(n \cdot \text{iterations})$

Stochastic Gradient Descent (SGD):

- Uses 1 sample per iteration (or mini-batch of size B)
- Per-iteration cost: $O(d)$ or $O(Bd)$
- Iterations needed: $O(1/\epsilon)$ for strongly convex, $O(1/\epsilon^2)$ for convex
- **Total sample accesses:** $O(\text{iterations})$

Sample Complexity: Detailed Analysis

For μ -strongly convex, L -smooth objectives:

Method	Iterations	Cost/Iter	Total Cost
GD	$O(\kappa \log(1/\epsilon))$	$O(nd)$	$O(\kappa nd \log(1/\epsilon))$
SGD	$O(1/(\mu\epsilon))$	$O(d)$	$O(d/(\mu\epsilon))$

where $\kappa = L/\mu$ is the condition number

Key observations:

- GD: logarithmic dependence on ϵ (fast convergence)
- SGD: linear dependence on $1/\epsilon$ (slower convergence per iteration)
- SGD wins when $n \gg 1/(\mu\epsilon)$ (large datasets!)
- Crossover point: $n \approx 1/(\mu\epsilon)$

Sample Complexity: When Does SGD Win?

Comparison for strongly convex case:

GD wins if: $\kappa n \log(1/\epsilon) < 1/(\mu\epsilon)$

- Small datasets (n small)
- Well-conditioned problems (κ small)
- High accuracy required (ϵ very small)

SGD wins if: $n \gg \frac{1}{\mu\epsilon\kappa \log(1/\epsilon)}$

- Large datasets (n large) - **typical in modern ML!**
- Moderate accuracy sufficient
- Can tolerate noisier updates

Example: $\epsilon = 0.01, \mu = 0.1, \kappa = 100$

- GD needs $\sim 460n$ sample accesses
- SGD needs ~ 100 sample accesses
- SGD wins decisively when $n > 100$!

Sample Complexity: Practical Implications

Modern machine learning regime:

- Large datasets: $n = 10^6$ to 10^9 (ImageNet, web-scale data)
- Moderate accuracy often sufficient: $\epsilon = 0.01$ or 0.001
- SGD-based methods dominate in practice

Additional SGD advantages:

- **Memory efficiency:** $O(d)$ vs $O(nd)$
- **Online learning:** can process streaming data
- **Generalization:** implicit regularization, better test performance
- **Hardware:** mini-batches fit in GPU memory

Trade-offs:

- + SGD: fewer total sample accesses, memory efficient
- SGD: more iterations, noisier, harder to tune
- + GD: stable, predictable convergence
- GD: expensive per iteration, not scalable to large n

Gradient Descent for High-Dimensional Linear Regression

Population Loss: For linear regression on data distribution

$$L(w) = \mathbb{E}_{(x,y) \sim \mathcal{D}}[(y - w^\top x)^2]$$

Gradient of population loss:

$$\begin{aligned}\nabla L(w) &= \mathbb{E}_{(x,y)}[\nabla_w(y - w^\top x)^2] \\ &= \mathbb{E}_{(x,y)}[-2(y - w^\top x)x] \\ &= -2\mathbb{E}[(y - w^\top x)x]\end{aligned}$$

Gradient Descent update rule:

$$w_{t+1} = w_t - \eta \nabla L(w_t)$$

where η is the learning rate (stepsize)

GD Iterates: Exact Form

Data model: Assume $y = w^* \top x + \xi$ where $\mathbb{E}[\xi|x] = 0$ (noise)

Define: Covariance matrix $H = \mathbb{E}[xx^\top]$

Gradient at w :

$$\begin{aligned}\nabla L(w) &= -2\mathbb{E}[(y - w^\top x)x] \\ &= -2\mathbb{E}[(w^{*\top} x + \xi - w^\top x)x] \\ &= -2\mathbb{E}[((w^* - w)^\top x)x] \\ &= -2H(w^* - w) = 2H(w - w^*)\end{aligned}$$

GD update becomes:

$$w_{t+1} = w_t - 2\eta H(w_t - w^*) = (I - 2\eta H)w_t + 2\eta Hw^*$$

GD Iterates: Closed Form

Recursive formula:

$$w_{t+1} - w^* = (I - 2\eta H)(w_t - w^*)$$

Unrolling the recursion:

$$\begin{aligned} w_t - w^* &= (I - 2\eta H)(w_{t-1} - w^*) \\ &= (I - 2\eta H)^2(w_{t-2} - w^*) \\ &= \dots \\ &= (I - 2\eta H)^t(w_0 - w^*) \end{aligned}$$

Exact form of iterates:

$$w_t = w^* + (I - 2\eta H)^t(w_0 - w^*)$$

Key insight: Convergence depends on eigenvalues of $(I - 2\eta H)$

Loss Function at Iterate w_t

Goal: Calculate $L(w_t) - L(w^*)$

Expand the loss:

$$\begin{aligned}L(w) &= \mathbb{E}[(y - w^\top x)^2] \\&= \mathbb{E}[(w^{*\top} x + \xi - w^\top x)^2] \\&= \mathbb{E}[((w^* - w)^\top x + \xi)^2] \\&= \mathbb{E}[(w^* - w)^\top x x^\top (w^* - w)] + \mathbb{E}[\xi^2] \\&= (w^* - w)^\top H(w^* - w) + \sigma^2\end{aligned}$$

where $\sigma^2 = \mathbb{E}[\xi^2]$ is the noise variance

Optimal loss:

$$L(w^*) = \sigma^2$$

Excess Loss and Convergence Rate

Excess loss at iteration t :

$$\begin{aligned} L(w_t) - L(w^*) &= (w^* - w_t)^\top H(w^* - w_t) \\ &= \|w^* - w_t\|_H^2 \end{aligned}$$

where $\|v\|_H^2 = v^\top H v$ is the Mahalanobis norm

Substitute the iterate formula:

$$\begin{aligned} L(w_t) - L(w^*) &= \|(I - 2\eta H)^t (w_0 - w^*)\|_H^2 \\ &= (w_0 - w^*)^\top [(I - 2\eta H)^t]^\top H [(I - 2\eta H)^t] (w_0 - w^*) \end{aligned}$$

Eigenvalue analysis: If H has eigenvalues $\lambda_1 \geq \dots \geq \lambda_d$:

- Convergence rate in direction v_i : $(1 - 2\eta\lambda_i)^{2t}$
- Fast convergence: large λ_i (high-variance directions)
- Slow convergence: small λ_i (low-variance directions)

Streaming SGD: Stochastic Version of GD

Constant-stepsize Streaming SGD:

In each iteration t , sample fresh (x_t, y_t) from data distribution:

$$w_t = w_{t-1} + \gamma \cdot (y_t - \langle w_{t-1}, x_t \rangle) \cdot x_t$$

- γ : constant stepsize
- $(y_t - \langle w_{t-1}, x_t \rangle) \cdot x_t$: stochastic gradient

Averaging the iterates:

- With constant stepsize, SGD oscillates around optimum
- Averaging reduces variance and gives better solution

$$\bar{w} = \frac{1}{K} \sum_{t=n-K}^{n-1} w_t.$$

High-Dimensional Linear Regression: Convergence Behavior

Key insight: In high dimensions, SGD/GD learn different features at different rates

Convergence rate depends on dimension:

- For linear regression with covariance matrix $H = \mathbb{E}[xx^\top]$
- H has eigenvalues $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_d$
- Dimensions with large eigenvalues: learn quickly
- Dimensions with small eigenvalues: learn slowly or not at all

Implications:

- **High-eigenvalue directions:** Converge in $O(\log(1/\epsilon))$ iterations
- **Low-eigenvalue directions:** May need $\gg d$ iterations
- GD/SGD provides **implicit regularization** by focusing on high-variance directions
- Similar to Ridge regression but without explicit λ tuning

Practical takeaway: GD/SGD naturally prioritizes important features in data

Simulation: Eigenvalue Decay $\lambda_k = 1/k$

Setup:

- Dimension $d = 20$
- Covariance matrix H with eigenvalues $\lambda_k = 1/k$ for $k = 1, \dots, d$
- Learning rate $\eta = 0.05$
- Run GD: $w_{t+1} = w_t - 2\eta H(w_t - w^*)$

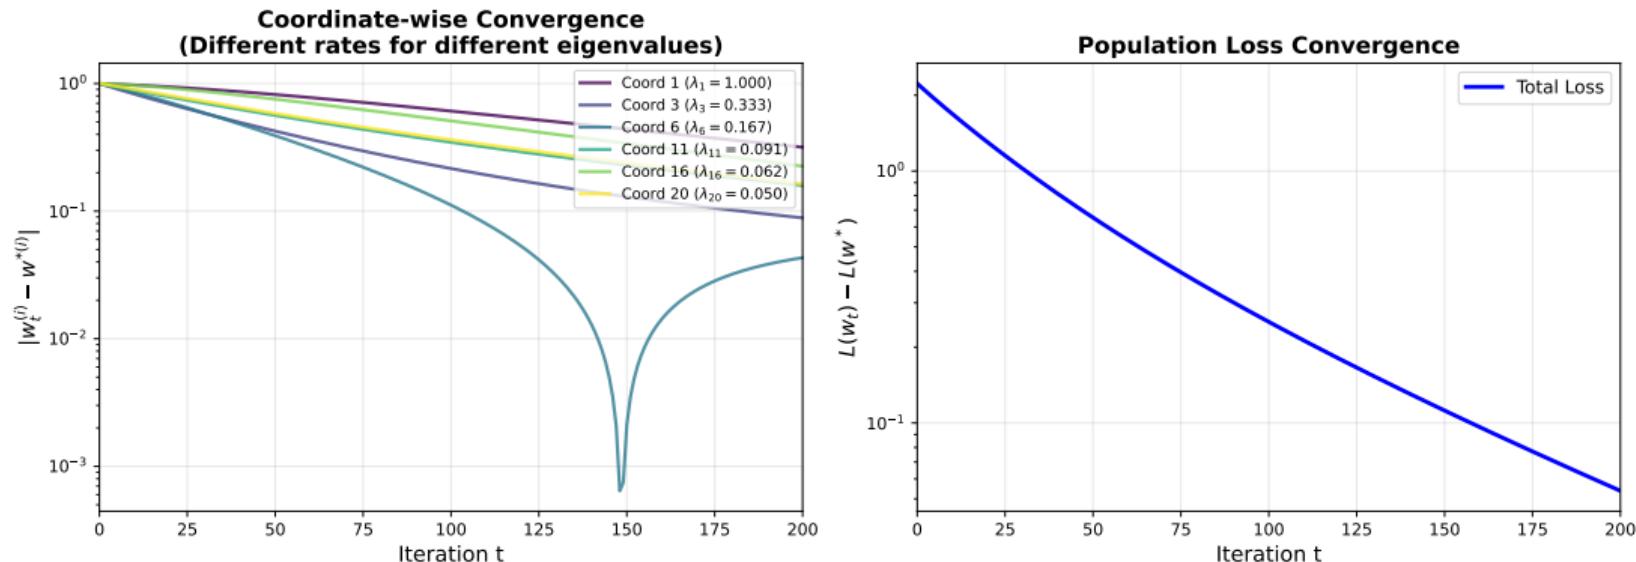
Theory predicts:

- Coordinate i converges at rate $(1 - 2\eta\lambda_i)^t$
- Large λ_i (first coordinates): fast convergence
- Small λ_i (last coordinates): slow convergence

Observations:

- First coordinate ($\lambda_1 = 1.0$): converges in ~ 20 iterations
- Middle coordinates ($\lambda_{10} \approx 0.1$): converges in ~ 100 iterations
- Last coordinates ($\lambda_{20} = 0.05$): still converging after 200 iterations

Coordinate-wise Convergence Dynamics



Key observations:

- Different coordinates converge at vastly different rates
- Convergence rate proportional to eigenvalue magnitude
- High-eigenvalue directions dominate early iterations

Neural Network Structure

Multi-layer network:

Input $x \rightarrow \text{Layer 1} \rightarrow \text{Layer 2} \rightarrow \dots \rightarrow \text{Layer } L \rightarrow \text{Output } f(x; \theta)$

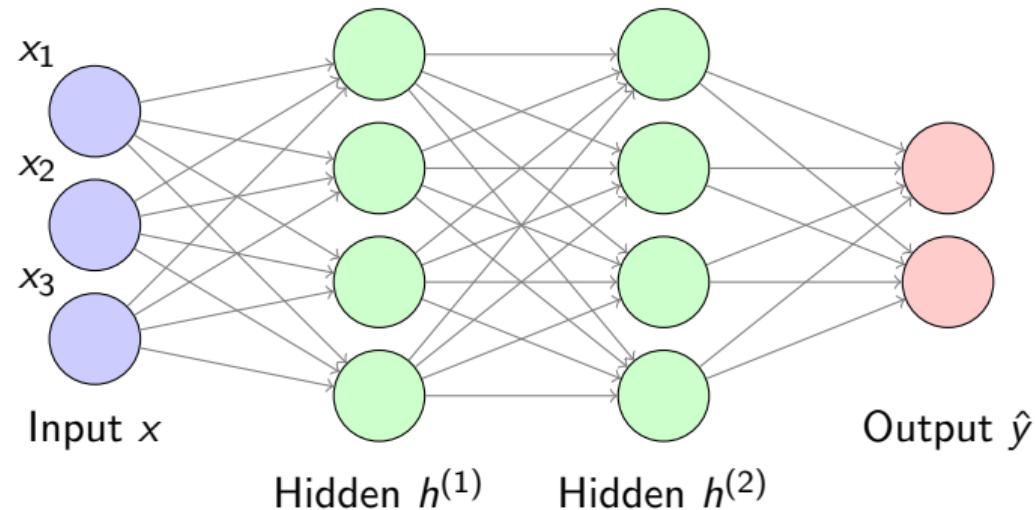
Layer computation:

$$h^{(\ell)} = \sigma(W^{(\ell)} h^{(\ell-1)} + b^{(\ell)})$$

where:

- $W^{(\ell)}$: weight matrix at layer ℓ
- $b^{(\ell)}$: bias vector
- σ : activation function (ReLU, sigmoid, tanh)
- $h^{(0)} = x$ (input)

Neural Network Visualization



Chain Rule for Gradients

Loss function: $L(\theta) = \ell(f(x; \theta), y)$

Goal: Compute $\nabla_{\theta} L$ for gradient descent

Chain rule: For nested functions $f(g(h(x)))$:

$$\frac{df}{dx} = \frac{df}{dg} \cdot \frac{dg}{dh} \cdot \frac{dh}{dx}$$

For neural networks:

- Forward pass: compute activations layer by layer
- Backward pass: compute gradients using chain rule
- Gradients flow backward through the network

Backpropagation Algorithm

Forward Pass (compute activations):

- 1: **for** $\ell = 1$ to L **do**
- 2: $z^{(\ell)} = W^{(\ell)} h^{(\ell-1)} + b^{(\ell)}$
- 3: $h^{(\ell)} = \sigma(z^{(\ell)})$
- 4: **end for**
- 5: Output: $f(x; \theta) = h^{(L)}$

Backward Pass (compute gradients):

- 1: Compute output gradient: $\delta^{(L)} = \nabla_{h^{(L)}} \ell(f(x; \theta), y)$
- 2: **for** $\ell = L$ down to 1 **do**
- 3: $\delta^{(\ell)} = \delta^{(\ell+1)} \odot \sigma'(z^{(\ell)})$ (*element-wise*)
- 4: $\nabla_{W^{(\ell)}} L = \delta^{(\ell)} (h^{(\ell-1)})^\top$
- 5: $\nabla_{b^{(\ell)}} L = \delta^{(\ell)}$
- 6: $\delta^{(\ell-1)} = (W^{(\ell)})^\top \delta^{(\ell)}$
- 7: **end for**

Backpropagation: Key Formula

Chain rule applied:

$$\delta^{(\ell)} = \frac{\partial L}{\partial z^{(\ell)}} = \frac{\partial L}{\partial z^{(\ell+1)}} \cdot \frac{\partial z^{(\ell+1)}}{\partial h^{(\ell)}} \cdot \frac{\partial h^{(\ell)}}{\partial z^{(\ell)}}$$

Intuition:

- $\delta^{(\ell)}$: sensitivity of loss to pre-activation at layer ℓ
- Propagates backward using chain rule
- Each layer: multiply by local derivatives

Efficiency:

- Reuses intermediate activations from forward pass
- Avoids recomputing derivatives multiple times
- Critical for training deep networks

Example: 2-Layer Network

Network:

$$h^{(1)} = \sigma(W^{(1)}x + b^{(1)})$$

$$f(x) = W^{(2)}h^{(1)} + b^{(2)}$$

Loss: $L = \frac{1}{2}\|f(x) - y\|^2$

Gradients:

Output layer:

$$\delta^{(2)} = f(x) - y$$

$$\nabla_{W^{(2)}} L = \delta^{(2)}(h^{(1)})^\top, \quad \nabla_{b^{(2)}} L = \delta^{(2)}$$

Hidden layer:

$$\delta^{(1)} = [(W^{(2)})^\top \delta^{(2)}] \odot \sigma'(W^{(1)}x + b^{(1)})$$

$$\nabla_{W^{(1)}} L = \delta^{(1)}x^\top, \quad \nabla_{b^{(1)}} L = \delta^{(1)}$$

SGD for Neural Network Training

Mini-batch SGD with backpropagation:

- 1: Sample mini-batch $\{(x_i, y_i)\}_{i \in \mathcal{I}}$
- 2: **for** each (x_i, y_i) **do**
- 3: Forward pass: compute $f(x_i; \theta)$
- 4: Backward pass: compute $\nabla_{\theta} \ell(f(x_i; \theta), y_i)$
- 5: **end for**
- 6: Average gradients: $g = \frac{1}{|\mathcal{I}|} \sum_{i \in \mathcal{I}} \nabla_{\theta} \ell(f(x_i; \theta), y_i)$
- 7: Update parameters: $\theta \leftarrow \theta - \eta \cdot g$

Computational Efficiency:

- Backpropagation: $O(\# \text{parameters})$ per sample
- Without backpropagation: would need $O(\# \text{parameters}^2)$
- **Critical for training large networks**

Summary: Gradient Descent

- ✓ Fundamental optimization algorithm for ML
- ✓ Convergence depends on function properties
 - Smooth: $O(1/t)$ for gradient norm
 - Smooth + Convex: $O(1/t)$ for function value
 - Smooth + Strongly convex: $O(\exp(-t/\kappa))$ exponential
- ✓ Requires computing full gradient: $O(n)$ per iteration
- ✗ Can get trapped at local minima or saddle points (non-convex)
- ✗ Expensive for large datasets

Summary: Stochastic Gradient Descent

- ✓ Uses mini-batches or single samples: $O(B)$ or $O(1)$ per iteration
- ✓ Expected update direction = gradient direction
- ✓ Can process streaming data
- ✓ Noise helps escape saddle points
- ✗ Has variance → needs careful stepsize selection
- ✗ Constant stepsize: oscillates around optimum
- ✗ Decaying stepsize: converges but slower per iteration

Summary: Applications

Linear Models:

- Regression, logistic: convex, well-understood
- SGD for linear regression: implicit regularization
- Competitive with ridge regression in many cases

Neural Networks:

- Backpropagation for efficient gradient computation
- SGD is the workhorse for deep learning
- Works empirically despite non-convexity

Practical Considerations:

- Stepsize selection is crucial (typically $\eta \leq 1/L$)
- Mini-batch size: trade-off between computation and variance
- For large datasets, SGD variants are essential

Key Takeaways

Why Numerical Optimization?

Modern ML problems are high-dimensional with complicated objectives. Gradient-based methods are the **only viable approach**.

Gradient Descent

Fundamental algorithm with strong theoretical guarantees for convex problems. Limitations in non-convex settings (local minima, saddle points).

Stochastic Gradient Descent

Scales to large datasets via mini-batching. Trade-off: faster iterations but noisier updates. Essential for modern deep learning.

Backpropagation

Efficient computation of gradients using chain rule. Enables training of deep neural networks.

Looking Forward

Topics beyond this lecture:

- **Momentum methods:** Accelerate convergence by using historical gradients
- **Adaptive learning rates:** AdaGrad, RMSProp, Adam
- **Variance reduction:** SVRG, SAGA for combining GD and SGD benefits
- **Second-order methods:** Newton's method, L-BFGS
- **Non-convex optimization theory:** Understanding why SGD works for deep learning

Thank you!